

1 Class-Conditional Densities for Binary Data

Question A:

$$p(x|y=c) = p(x_1|y=c)p(x_2|x_1, y=c) \cdots p(x_D|x_1, \dots, x_{D-1}, y=c) = \prod_{j=1}^D \theta_{x_j c}$$

Since all the D features are binary, $x_j \in \{0, 1\}$, for each class of C we need 2^{j-1} parameters for $p(x_j|x_1, \dots, x_{j-1}, y=c)$. In total $\sum_{j=1}^D 2^{j-1} = O(2^D)$. Therefore, we need $O(C \times 2^D)$ parameters for C classes.

Question B: Without factorization, since all the D features are binary, $x_j \in \{0, 1\}$, and C classes for y , in total we need $O(C \times 2^D)$, which is the same as that with factorization.

Question C: For a small N , Naive Bayes is likely to give lower test set error. This is because full models are more likely to overfit with a small N .

Question D: For a large N , full models are likely to give lower test set error. This is because Naive Bayes is simple and is likely to underfit with a large N , while full models have more parameters and will do better.

Question E: For Naive Bayes,

$$p(y|x) = \frac{p(x, y)}{p(x)} = \frac{p(y)}{p(x)} \prod_d p(x^d|y) = \frac{p(y)}{\sum_{i=1}^C p(x|y=c_i)} \prod_d p(x^d|y)$$

Since we assumed a uniform class prior $p(y)$, $O(p(y)) = O(1)$, and computation complexity of $p(y|x) = O(CD)$.

For full model,

$$p(y|x) = \frac{p(x, y)}{p(x)} = \frac{p(y)}{p(x)} p(x|y)$$

, and computation complexity of D -dimensional vector is $O(D)$. Different from Naive Bayes, full model doesn't take into account every $y=c$, but their overall probability. Therefore, computation complexity of $p(y|x) = O(D)$.

2 Sequence Prediction

Question A:

```
#####
Running Code For Question 2A
#####

File #0:
Emission Sequence      Max Probability State Sequence
#####
25421                  31033
01232367534          22222100310
5452674261527433     1031003103222222
7226213164512267255  1310331000033100310
0247120602352051010255241 222222222222222222222103

File #1:
Emission Sequence      Max Probability State Sequence
#####
77550                  22222
7224523677            2222221000
505767442426747       221100003310031
72134131645536112267  1031031000031033100
4733667771450051060253041 221100003322223103222223

File #2:
Emission Sequence      Max Probability State Sequence
#####
00622                  11111
4687981156            2100202111
815833657775062       021011111111111
21310222515963505015  02020111111111111021
6503199452571274006320025 1110202111111102021110211

File #3:
Emission Sequence      Max Probability State Sequence
#####
13661                  00021
2102213421            3131310213
166066262165133      133333133133100
53164662112162634156 20000021313131002133
1523541005123230226306256 13100213313313313133133

File #4:
Emission Sequence      Max Probability State Sequence
#####
23664                  01124
3630535602            0111201112
350201162150142       01124012441112
00214005402015146362  11201112412444011112
2111266524665143562534450 2012012424124011112411124

File #5:
Emission Sequence      Max Probability State Sequence
#####
68535                  10111
4546566636            1111111111
638436858101213       110111010000011
13240338308444514688  00010000000111111100
0111664434441382533632626 2111111111111100111110101
```

Question B: The results using Forward algorithm are as follows:

```
#####
Running Code For Question 2B1
#####

File #0:
Emission Sequence      Probability of Emitting Sequence
#####
25421                  4.537e-05
01232367534          1.620e-11
5452674261527433     4.348e-15
7226213164512267255  4.739e-18
0247120602352051010255241 9.365e-24

File #1:
Emission Sequence      Probability of Emitting Sequence
#####
77550                  1.181e-04
7224523677            2.033e-09
505767442426747       2.477e-13
72134131645536112267  8.871e-20
4733667771450051060253041 3.740e-24

File #2:
Emission Sequence      Probability of Emitting Sequence
#####
00622                  2.088e-05
4687981156            5.181e-11
815833657775062       3.315e-15
21310222515963505015  5.126e-20
6503199452571274006320025 1.297e-25

File #3:
Emission Sequence      Probability of Emitting Sequence
#####
13661                  1.732e-04
2102213421            8.285e-09
166066262165133      1.642e-12
53164662112162634156  1.063e-16
1523541005123230226306256 4.535e-22

File #4:
Emission Sequence      Probability of Emitting Sequence
#####
23664                  1.141e-04
3630535602            4.326e-09
350201162150142       9.793e-14
00214005402015146362  4.740e-18
2111266524665143562534450 5.618e-22

File #5:
Emission Sequence      Probability of Emitting Sequence
#####
68535                  1.322e-05
4546566636            2.867e-09
638436858101213       4.323e-14
13240338308444514688  4.620e-18
0111664434441382533632626 1.440e-22
```

The results using Backward algorithm are as follows:

```

=====
Running Code For Question 2Bii
=====

File #0:
Emission Sequence      Probability of Emitting Sequence
=====
25421                  4.537e-05
01232367534           1.620e-11
5452674261527433      4.348e-15
7226213164512267255   4.739e-18
0247120602352051010255241 9.365e-24

File #1:
Emission Sequence      Probability of Emitting Sequence
=====
77550                  1.181e-04
7224523677            2.033e-09
505767442426747       2.477e-13
72134131645536112267   8.871e-20
4733667771450051060253041 3.740e-24

File #2:
Emission Sequence      Probability of Emitting Sequence
=====
60622                  2.088e-05
4687981156             5.181e-11
815833657775062        3.315e-15
21310222515963505015   5.126e-20
6503199452571274006320025 1.297e-25

File #3:
Emission Sequence      Probability of Emitting Sequence
=====
13661                  1.732e-04
2102213421             8.285e-09
166066262165133        1.642e-12
53164662112162634156   1.063e-16
1523541005123230226306256 4.535e-22

File #4:
Emission Sequence      Probability of Emitting Sequence
=====
23664                  1.141e-04
3630535602             4.326e-09
350201162150142        9.793e-14
00214005402015146362   4.740e-18
2111266524665143562534450 5.618e-22

File #5:
Emission Sequence      Probability of Emitting Sequence
=====
68535                  1.322e-05
4546566636             2.867e-09
638436858101213        4.322e-14
13240338308444514688   4.620e-18
0111664434441382533632626 1.440e-22
    
```

Question C:

```

=====
Running Code For Question 2C
=====

Transition Matrix:
=====
2.833e-01  4.714e-01  1.310e-01  1.143e-01
2.321e-01  3.810e-01  2.940e-01  9.284e-02
1.040e-01  9.760e-02  3.696e-01  4.288e-01
1.883e-01  9.903e-02  3.052e-01  4.075e-01

Observation Matrix:
=====
1.486e-01  2.288e-01  1.533e-01  1.179e-01  4.717e-02  5.189e-02  2.830e-02  1.297e-01  9.198e-02  2.358e-03
1.062e-01  9.653e-03  1.931e-02  3.089e-02  1.699e-01  4.633e-02  1.409e-01  2.394e-01  1.371e-01  1.004e-01
1.194e-01  4.299e-02  6.529e-02  9.076e-02  1.768e-01  2.022e-01  4.618e-02  5.096e-02  7.803e-02  1.274e-01
1.694e-01  3.871e-02  1.468e-01  1.823e-01  4.839e-02  6.290e-02  9.032e-02  2.581e-02  2.161e-01  1.935e-02
    
```

Question D:

```
#####
Running Code For Question 2D
#####

test case: N_iters=0
Transition Matrix:
#####
2.470e-01  6.958e-02  3.064e-01  3.770e-01
1.846e-01  3.635e-01  2.114e-01  2.404e-01
2.902e-01  1.522e-01  1.328e-01  4.248e-01
3.266e-01  2.547e-01  5.671e-02  3.620e-01

Observation Matrix:
#####
9.093e-02  8.443e-02  1.399e-01  1.539e-01  1.214e-01  6.519e-02  8.377e-02  1.268e-01  9.067e-02  4.306e-02
4.669e-02  2.347e-01  1.511e-02  5.726e-03  1.281e-01  4.333e-02  5.775e-02  1.542e-01  1.871e-01  1.273e-01
1.797e-01  1.101e-01  6.323e-02  1.346e-01  3.886e-02  9.526e-02  1.189e-01  2.195e-02  8.021e-02  1.572e-01
1.381e-01  2.190e-01  6.830e-02  1.514e-01  1.058e-01  1.222e-02  1.477e-01  1.160e-01  7.535e-03  3.392e-02

test case: N_iters=1
Transition Matrix:
#####
2.702e-01  7.208e-02  3.259e-01  3.318e-01
2.003e-01  3.712e-01  2.216e-01  2.069e-01
3.229e-01  1.585e-01  1.439e-01  3.747e-01
3.561e-01  2.646e-01  6.016e-02  3.192e-01

Observation Matrix:
#####
9.941e-02  3.250e-02  1.594e-01  1.280e-01  1.222e-01  1.215e-01  5.508e-02  1.076e-01  1.409e-01  3.342e-02
5.500e-02  9.644e-02  1.872e-02  4.955e-03  1.382e-01  8.785e-02  4.052e-02  1.451e-01  3.061e-01  1.072e-01
1.987e-01  4.211e-02  7.324e-02  1.137e-01  4.002e-02  1.815e-01  8.075e-02  1.931e-02  1.253e-01  1.253e-01
1.880e-01  1.038e-01  9.799e-02  1.553e-01  1.338e-01  2.815e-02  1.223e-01  1.229e-01  1.453e-02  3.325e-02

Transition Matrix:
#####
8.029e-04  5.274e-01  8.376e-05  4.717e-01
1.856e-03  6.830e-01  2.922e-01  2.303e-02
6.214e-01  8.278e-13  3.747e-01  3.940e-03
2.051e-02  7.025e-01  2.042e-02  2.566e-01

Observation Matrix:
#####
1.577e-01  4.863e-02  1.835e-01  4.863e-02  2.072e-01  3.686e-21  6.874e-02  2.040e-02  2.653e-01  1.458e-35
1.120e-01  5.788e-02  1.132e-01  1.021e-01  1.309e-01  9.003e-02  8.507e-14  1.870e-01  1.437e-01  6.304e-02
9.063e-02  7.887e-02  1.014e-16  1.937e-01  6.975e-02  2.191e-01  1.481e-01  9.426e-17  5.298e-02  1.468e-01
3.079e-01  1.320e-01  9.116e-02  3.171e-02  5.488e-03  1.068e-06  2.835e-01  6.466e-02  8.364e-02  4.515e-07
```

Question E: The transition and emission matrices from 2C and 2D are different, supervised matrices are more uniform, while unsupervised are sparse and nonuniform. 2C (supervised HMM) provides a more accurate representation. This is because for unsupervised 2D, we simply random the initial matrices, which cannot truly reflect Ron's moods. To improve the unsupervised learning data, we might need to initialize the transition and emission matrices A & O more precisely, such as using probability distributions from supervised matrices.

Question F:

```
#####
Running Code For Question 2F
#####

File #0:
Generated Emission
#####
45274511250275204412
05407242224726770147
72422713473031327660
42427312074107504324
65147177214245322422

File #1:
Generated Emission
#####
40035770104447412207
45516552177027677212
0431742222000147762
35746073172105617607
05107555006525007447

File #2:
Generated Emission
#####
17726563315712157156
90030125025149435211
62975620039788065245
45692563828145922702
52899330755675696052

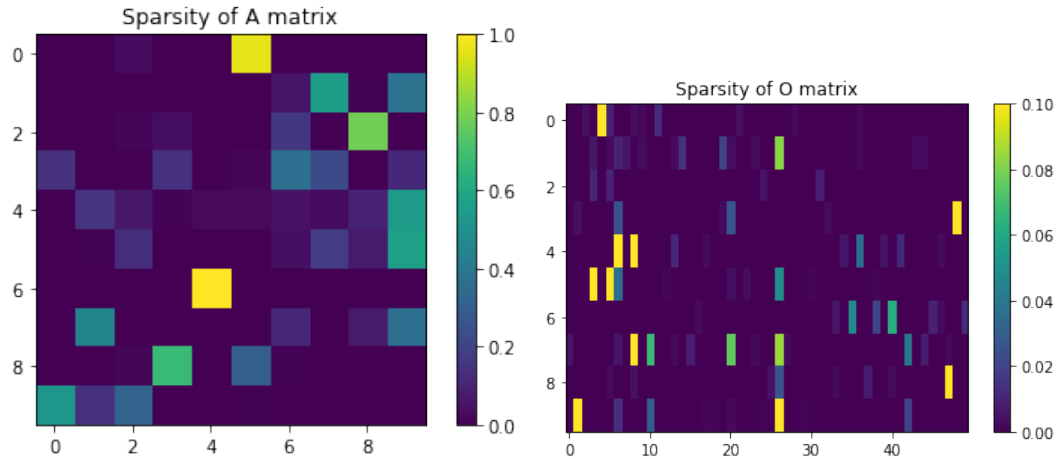
File #3:
Generated Emission
#####
03341033031362462112
15126633304022436420
62165152146652334362
12252065242314333013
21214502130111364421

File #4:
Generated Emission
#####
34202331563531320223
16566131513143621520
43335062640460660662
62431603632366426636
16354030615436116252

File #5:
Generated Emission
#####
32511566801635284208
81164464364484665615
04666240667481480436
60366565318454544314
45422451038680813715
```

Question G:

The trained A and O matrices are both sparse, with most elements near 0. The sparsity of transition matrix A means that for each state, there might be very few states to transit to, and the sparsity of observation matrix O means that for each state, there might be very few observations to belong to.



Question H: As the number of hidden states is increased, sample emission sentences from the HMM becomes more coherent and smooth. When there is only one hidden state, then there is no more transitions, meaning that observation words are just randomly picked from the dataset. Allowing more hidden states will increase the training data likelihood, but when we have too many hidden states for the fixed observation set, this will lead to overfitting and may increase the test error.

Question I: As shown in the figure, I think state 7 is semantically meaningful. This state is filled with



law-related words, such as 'law', 'court', 'representative', 'justice', 'rules'. I think this state represent all law-related words in the dataset, which is distinct from other states.

HMM.py

March 5, 2020

```
[ ]: #####
# CS/CNS/EE 155 2018
# Problem Set 6
#
# Author:      Andrew Kang
# Description: Set 6 skeleton code
#####

# You can use this (optional) skeleton code to complete the HMM
# implementation of set 5. Once each part is implemented, you can simply
# execute the related problem scripts (e.g. run 'python 2G.py') to quickly
# see the results from your code.
#
# Some pointers to get you started:
#
# - Choose your notation carefully and consistently! Readable
#   notation will make all the difference in the time it takes you
#   to implement this class, as well as how difficult it is to debug.
#
# - Read the documentation in this file! Make sure you know what
#   is expected from each function and what each variable is.
#
# - Any reference to "the  $(i, j)$ -th" element of a matrix  $T$  means that
#   you should use  $T[i][j]$ .
#
# - Note that in our solution code, no NumPy was used. That is, there
#   are no fancy tricks here, just basic coding. If you understand HMMs
#   to a thorough extent, the rest of this implementation should come
#   naturally. However, if you'd like to use NumPy, feel free to.
#
# - Take one step at a time! Move onto the next algorithm to implement
#   only if you're absolutely sure that all previous algorithms are
#   correct. We are providing you waypoints for this reason.
#
# To get started, just fill in code where indicated. Best of luck!

import random
```

```

class HiddenMarkovModel:
    '''
    Class implementation of Hidden Markov Models.
    '''

    def __init__(self, A, O):
        '''
        Initializes an HMM. Assumes the following:
        - States and observations are integers starting from 0.
        - There is a start state (see notes on A_start below). There
          is no integer associated with the start state, only
          probabilities in the vector A_start.
        - There is no end state.

        Arguments:
        A:          Transition matrix with dimensions L x L.
                   The (i, j)th element is the probability of
                   transitioning from state i to state j. Note that
                   this does not include the starting probabilities.

        O:          Observation matrix with dimensions L x D.
                   The (i, j)th element is the probability of
                   emitting observation j given state i.

        Parameters:
        L:          Number of states. word of tags: N, V, adj

        D:          Number of observations. fish, sleep

        A:          The transition matrix. L x L

        O:          The observation matrix. L x D

        A_start:    Starting transition probabilities. The ith element
                   is the probability of transitioning from the start
                   state to state i. For simplicity, we assume that
                   this distribution is uniform.
        '''

        self.L = len(A)
        self.D = len(O[0])
        self.A = A
        self.O = O
        self.A_start = [1. / self.L for _ in range(self.L)]

```



```

def viterbi(self, x):
    '''
    Uses the Viterbi algorithm to find the max probability state
    sequence corresponding to a given input sequence.

    Arguments:
        x:          Input sequence in the form of a list of length M,
                    consisting of integers ranging from 0 to D - 1.

    Returns:
        max_seq:    State sequence corresponding to x with the highest
                    probability.
    '''

    M = len(x)      # Length of sequence.

    # The (i, j)-th elements of probs and seqs are the max probability
    # of the prefix of length i ending in state j and the prefix
    # that gives this probability, respectively.
    #
    # For instance, probs[1][0] is the probability of the prefix of
    # length 1 ending in state 0.
    probs = [[0. for _ in range(self.L)] for _ in range(M + 1)] #
    ↪ probability
    seqs = [['' for _ in range(self.L)] for _ in range(M + 1)] # viterbi
    ↪

    # L # of states/tags, D: # of observations/ , M: sequence , A: L x L,
    ↪ 0: L x D

    for i in range(self.L): # initialize 1st state,      probs[1][i] A_start
    ↪ * 0, probs[0][0,...,0]
        probs[1][i] = self.A_start[i] * self.O[i][x[0]]

    for seq in range(1, M): # sequence
        for curr in range(self.L): # states/tags
            max_value = 0
            max_idx = 0
            for last in range(self.L): # state
                if (probs[seq][last] * self.A[last][curr] * self.
    ↪ 0[curr][x[seq]] >= max_value):
                    max_value = probs[seq][last] * self.A[last][curr] *
    ↪ self.O[curr][x[seq]] #
                    max_idx = last
            probs[seq + 1][curr] = max_value # state probs
            seqs[seq + 1][curr] = max_idx

```

```

max_seq_rev = []

max_value = max(probs[M])
max_idx = probs[M].index(max_value) #
↪ state max_value( max_prob) max_idx

max_seq_rev.append( str(max_idx) )

for i in range(M, 1, -1):
    max_seq_rev.append( str(seqs[i][max_idx]) )
    max_idx = seqs[i][max_idx]

max_seq = max_seq_rev[::-1]

return "".join(max_seq)

def forward(self, x, normalize=False):
    '''
    Uses the forward algorithm to calculate the alpha probability
    vectors corresponding to a given input sequence.

    Arguments:
        x:          Input sequence in the form of a list of length M,
                    consisting of integers ranging from 0 to D - 1.

        normalize:  Whether to normalize each set of  $\alpha_j(i)$  vectors
                    at each  $i$ . This is useful to avoid underflow in
                    unsupervised learning.

    Returns:
        alphas:      Vector of alphas.

                    The  $(i, j)$ -th element of alphas is  $\alpha_j(i)$ ,
                    i.e. the probability of observing prefix  $x^1:i$ 
                    and state  $y^i = j$ .

                    e.g. alphas[1][0] corresponds to the probability
                    of observing  $x^1:1$ , i.e. the first observation,
                    given that  $y^1 = 0$ , i.e. the first state is 0.
    '''

    M = len(x)          # Length of sequence.
    alphas = [[0. for _ in range(self.L)] for _ in range(M + 1)]

    for i in range(self.L):

```

```

        alphas[1][i] = self.A_start[i] * self.O[i][x[0]]

    for seq in range(1, M): # sequence
        for curr in range(self.L): # states/tags
            sum_value = 0
            for last in range(self.L): # state
                sum_value += alphas[seq][last] * self.A[last][curr] * self.
→O[curr][x[seq]] # Viterbi replaces sum with max

            alphas[seq + 1][curr] = sum_value

    if normalize:
        sum_alpha = sum(alphas[seq + 1])
        for curr in range(self.L):
            alphas[seq + 1][curr] /= sum_alpha

    return alphas

def backward(self, x, normalize=False):
    """
    Uses the backward algorithm to calculate the beta probability
    vectors corresponding to a given input sequence.

    Arguments:
        x: Input sequence in the form of a list of length M,
            consisting of integers ranging from 0 to D - 1.

        normalize: Whether to normalize each set of beta_j(i) vectors
            at each i. This is useful to avoid underflow in
            unsupervised learning.

    Returns:
        betas: Vector of betas.

        The (i, j)-th element of betas is beta_j(i), i.e.
        the probability of observing prefix  $x^{(i+1):M}$  and
        state  $y^i = j$ .

        e.g. betas[M][0] corresponds to the probability
        of observing  $x^{M+1:M}$ , i.e. no observations,
        given that  $y^M = 0$ , i.e. the last state is 0.
    """

    M = len(x) # Length of sequence.
    betas = [[0. for _ in range(self.L)] for _ in range(M + 1)]

```

```

for i in range(self.L):
    betas[-1][i] = 1 # PPT 74,  $\beta(M) = 1$ 

for seq in range(-1, -M-1, -1): # sequence
    for curr in range(self.L): # states/tags
        sum_value = 0
        for nxt in range(self.L): # state
            if seq != -M:
                sum_value += betas[seq][nxt] * self.A[curr][nxt] * self.
↪ 0[nxt][x[seq]] # PPT 74,  $A_{\{z,j\}} A_{\{j,z\}}$ 
            else:
                sum_value += betas[seq][nxt] * self.A_start[nxt] * self.
↪ 0[nxt][x[seq]]

        betas[seq - 1][curr] = sum_value

    if normalize:
        sum_beta = sum(betas[seq - 1])
        for curr in range(self.L):
            betas[seq - 1][curr] /= sum_beta

return betas

def supervised_learning(self, X, Y):
    """
    Trains the HMM using the Maximum Likelihood closed form solutions
    for the transition and observation matrices on a labeled
    dataset (X, Y). Note that this method does not return anything, but
    instead updates the attributes of the HMM object.

    Arguments:
        X: A dataset consisting of input sequences in the form
            of lists of variable length, consisting of integers
            ranging from 0 to D - 1. In other words, a list of
            lists.

        Y: A dataset consisting of state sequences in the form
            of lists of variable length, consisting of integers
            ranging from 0 to L - 1. In other words, a list of
            lists.

            Note that the elements in X line up with those in Y.
    """
    # Calculate each element of A using the M-step formulas.

```

```

A_count = [[0. for i in range(self.L)] for j in range(self.L)]
A_sum = [0. for i in range(self.L)]

# For each input sequence:
for y in Y: # Y is a list of lists of length L, so y should be a list,
→ each y[i] is a mood tag
    for i in range(len(y) - 1): # A is calculated by
        A_count[ y[i] ][ y[i + 1] ] += 1
        A_sum[y[i]] += 1

for curr in range(self.L): # to normalize the A matrix
    for nxt in range(self.L):
        self.A[curr][nxt] = A_count[curr][nxt] / A_sum[curr]

# Calculate each element of O using the M-step formulas.

O_count = [[0. for i in range(self.D)] for j in range(self.L)] # O = L
→ x D, (i, j) is the probability of emitting observation j given state i.
O_sum = [0. for i in range(self.L)]

for x, y in zip(X, Y):
    for i in range(len(y)):
        O_count[ y[i] ][ x[i] ] += 1
        O_sum[ y[i] ] += 1

for curr in range(self.L):
    for nxt in range(self.D):
        self.O[curr][nxt] = O_count[curr][nxt] / O_sum[curr]

pass

def unsupervised_learning(self, X, N_iters):
    """
    Trains the HMM using the Baum-Welch algorithm on an unlabeled
    dataset X. Note that this method does not return anything, but
    instead updates the attributes of the HMM object.

    Arguments:
        X:          A dataset consisting of input sequences in the form
                    of lists of length M, consisting of integers ranging
                    from 0 to D - 1. In other words, a list of lists.

        N_iters:    The number of iterations to train on.
    """

```

```

for iteration in range(N_iters):

    A_count = [[0. for i in range(self.L)] for j in range(self.L)]
    A_sum = [0. for i in range(self.L)]
    O_count = [[0. for i in range(self.D)] for j in range(self.L)]
    O_sum = [0. for i in range(self.L)]

    # for each list
    for x in X:
        M = len(x)          # Length of sequence.
        # expectation step: given A & O matrix, predict probs of y's
        →for each training x
            # use forward-backward algorithm, alpha & beta [1,...,M]
            alphas = self.forward(x, normalize=True) # (M+1)xL, (i, j) is
            →alpha_j(i), probability of observing prefix x1:i and state yi = j.
            betas = self.backward(x, normalize=True)

            Marginals = [0. for _ in range(self.L)]
            for seq in range(1, M + 1): # for each prefix x1:i
                for curr in range(self.L): # states/tags
                    Marginals[curr] = alphas[seq][curr] * betas[seq][curr]

            Marginals_sum = sum(Marginals)
            for curr in range(self.L):
                Marginals[curr] /= Marginals_sum # normalized P(yi
            →(curr) | x), 0

            # Maximization Step: Use y's to estimate new (A,O)
            O_count[curr][x[seq - 1]] += Marginals[curr] #
            →seq 1 , x[] 0

            O_sum[curr] += Marginals[curr]
            if seq != M: # A is calculated by
                A_sum[curr] += Marginals[curr] # for unsupervised
            →learning, there is no tag, A_count

            # P(yi, yi+1 | x) A
            for seq in range(1, M):
                A_update = [[0. for _ in range(self.L)] for _ in range(self.
            →L)]

                for curr in range(self.L):
                    for nxt in range(self.L):
                        A_update[curr][nxt] = alphas[seq][curr] * self.
            →A[curr][nxt] * self.O[nxt][x[seq]] * betas[seq + 1][nxt]
                        # seq 1 , x[] 0

```

```

        A_update_sum = sum( [sum(A_update[i]) for i in
→range(len(A_update)) ] )

        for curr in range(self.L):
            for nxt in range(self.L):
                A_update[curr][nxt] /= A_update_sum # normalized
→P(yi, yi+1 | x)

        for curr in range(self.L):
            for nxt in range(self.L):
                A_count[curr][nxt] += A_update[curr][nxt]

        for curr in range(self.L): # to normalize the A & O matrix
            for nxt in range(self.L):
                self.A[curr][nxt] = A_count[curr][nxt] / A_sum[curr]

        for curr in range(self.L):
            for nxt in range(self.D):
                self.O[curr][nxt] = O_count[curr][nxt] / O_sum[curr]
pass

def get_data_with_distribute(self, dist): #
    r = random.random()
    for i, p in enumerate(dist):
        if r < p:
            return i
        r -= p

def generate_emission(self, M):
    """
    Generates an emission of length M, assuming that the starting state
    is chosen uniformly at random.

    Arguments:
        M:          Length of the emission to generate.

    Returns:
        emission:    The randomly generated emission as a list.

        states:      The randomly generated states as a list.
    """

    emission = []
    states = []

```

```

        y_start = random.randint(0, self.L - 1) # starting state is chosen
        ↪ uniformly at random
        states.append(y_start)

    for i in range(M):

        # Generate observation/emission x
        #idx_random = random.randint(0, self.D - 1)
        #max_value = max(self.O[ states[i] ])
        #x_index = self.O[ states[i] ].index(max_value)
        x_index = self.get_data_with_distribute(self.O[ states[i] ])
        emission.append(x_index)

        # Generate next state y.
        #idx_random = random.randint(0, self.L - 1)
        y_index = self.get_data_with_distribute(self.A[ states[i] ])
        states.append(y_index)

    return emission, states[:-1]

def probability_alphas(self, x):
    """
    Finds the maximum probability of a given input sequence using
    the forward algorithm.

    Arguments:
        x:          Input sequence in the form of a list of length M,
                    consisting of integers ranging from 0 to D - 1.

    Returns:
        prob:       Total probability that x can occur.
    """

    # Calculate alpha vectors.
    alphas = self.forward(x)

    # alpha_j(M) gives the probability that the state sequence ends
    # in j. Summing this value over all possible states j gives the
    # total probability of x paired with any state sequence, i.e.
    # the probability of x.
    prob = sum(alphas[-1])
    return prob

def probability_betas(self, x):
    """

```


Finds the maximum probability of a given input sequence using the backward algorithm.

Arguments:

x: Input sequence in the form of a list of length M, consisting of integers ranging from 0 to D - 1.

Returns:

prob: Total probability that x can occur.

'''

```
betas = self.backward(x)
```

```
# beta_j(1) gives the probability that the state sequence starts  
# with j. Summing this, multiplied by the starting transition  
# probability and the observation probability, over all states  
# gives the total probability of x paired with any state  
# sequence, i.e. the probability of x.
```

```
prob = sum([betas[1][j] * self.A_start[j] * self.O[j][x[0]] \  
            for j in range(self.L)])
```

```
return prob
```

```
def supervised_HMM(X, Y):
```

```
'''
```

Helper function to train a supervised HMM. The function determines the number of unique states and observations in the given data, initializes the transition and observation matrices, creates the HMM, and then runs the training function for supervised learning.

Arguments:

X: A dataset consisting of input sequences in the form of lists of variable length, consisting of integers ranging from 0 to D - 1. In other words, a list of lists.

Y: A dataset consisting of state sequences in the form of lists of variable length, consisting of integers ranging from 0 to L - 1. In other words, a list of lists. Note that the elements in X line up with those in Y.

```
'''
```

```
# Make a set of observations.
```

```
observations = set()
```

```
for x in X:
```

```
    observations |= set(x)
```

```
# Make a set of states.
```

```

states = set()
for y in Y:
    states |= set(y)

# Compute L and D.
L = len(states)
D = len(observations)

# Randomly initialize and normalize matrix A.
A = [[random.random() for i in range(L)] for j in range(L)]

for i in range(len(A)):
    norm = sum(A[i])
    for j in range(len(A[i])):
        A[i][j] /= norm

# Randomly initialize and normalize matrix O.
O = [[random.random() for i in range(D)] for j in range(L)]

for i in range(len(O)):
    norm = sum(O[i])
    for j in range(len(O[i])):
        O[i][j] /= norm

# Train an HMM with labeled data.
HMM = HiddenMarkovModel(A, O)
HMM.supervised_learning(X, Y)

return HMM

def unsupervised_HMM(X, n_states, N_iters):
    """
    Helper function to train an unsupervised HMM. The function determines the
    number of unique observations in the given data, initializes
    the transition and observation matrices, creates the HMM, and then runs
    the training function for unsupervised learning.

    Arguments:
        X: A dataset consisting of input sequences in the form
            of lists of variable length, consisting of integers
            ranging from 0 to D - 1. In other words, a list of lists.

        n_states: Number of hidden states to use in training.

        N_iters: The number of iterations to train on.
    """

```

```

# Make a set of observations.
observations = set()
for x in X:
    observations |= set(x)

# Compute L and D.
L = n_states
D = len(observations)

# Randomly initialize and normalize matrix A.
random.seed(2020)
A = [[random.random() for i in range(L)] for j in range(L)]

for i in range(len(A)):
    norm = sum(A[i])
    for j in range(len(A[i])):
        A[i][j] /= norm

# Randomly initialize and normalize matrix O.
random.seed(155)
O = [[random.random() for i in range(D)] for j in range(L)]

for i in range(len(O)):
    norm = sum(O[i])
    for j in range(len(O[i])):
        O[i][j] /= norm

# Train an HMM with unlabeled data.
HMM = HiddenMarkovModel(A, O)
HMM.unsupervised_learning(X, N_iters)

return HMM

```

2_notebook

March 5, 2020

1 Problem 2

In this Jupyter notebook, we visualize how HMMs work. This visualization corresponds to problem 2 in set 6.

Assuming your HMM module is complete and saved at the correct location, you can simply run all cells in the notebook without modification.

```
[1]: import os
import numpy as np
from IPython.display import HTML

from HMM import unsupervised_HMM
from HMM_helper import (
    text_to_wordcloud,
    states_to_wordclouds,
    parse_observations,
    sample_sentence,
    visualize_sparsities,
    animate_emission
)
```

1.1 Visualization of the dataset

We will be using the Constitution as our dataset. First, we visualize the entirety of the Constitution as a wordcloud:



```
[2]: text = open(os.path.join(os.getcwd(), 'data/constitution.txt')).read()
wordcloud = text_to_wordcloud(text, title='Constitution')
```

Now we train an HMM on our dataset. We use 10 hidden states and train over 100 iterations:

```
obs, obs_map = parse_observations(text)
hmm8 = unsupervised_HMM(obs, 10, 100)
```

We can visualize the sparsities of the A and O matrices by treating the matrix entries as intensity values and showing them as images. What patterns do you notice?

```
visualize_sparsities(hmm8, 0_max_cols=50)
```



1.4 Generating a sample sentence

As you have already seen, an HMM can be used to generate sample sequences based on the given dataset. Run the cell below to show a sample sentence based on the Constitution.

```
[5]: print('Sample Sentence:\n=====')
      print(sample_sentence(hmm8, obs_map, n_words=25))
```

Sample Sentence:

=====

War office the concurrence necessary and unless in state obliged to exports
thirds be public thereby the elections states and expiration of an and
counterfeiting..

1.5 Part H: Using varying numbers of hidden states

Using different numbers of hidden states can lead to different behaviours in the HMMs. Below, we train several HMMs with 1, 2, 4, and 16 hidden states, respectively. What do you notice about their emissions? How do these emissions compare to the emission above?

```
[6]: hmm1 = unsupervised_HMM(obs, 1, 100)
      print('\nSample Sentence:\n=====')
      print(sample_sentence(hmm1, obs_map, n_words=25))
```

Sample Sentence:

=====

Or justice the shall state the shall emit court effect together any pay shall
the and bound of states originated such the member capitation congress...

```
[7]: hmm2 = unsupervised_HMM(obs, 2, 100)
      print('\nSample Sentence:\n=====')
      print(sample_sentence(hmm2, obs_map, n_words=25))
```

Sample Sentence:

=====

Legislature tonnage and and president and electors representative deprived
effect or section and number between to that and shall the senate before by the
whom...

```
[8]: hmm4 = unsupervised_HMM(obs, 4, 100)
      print('\nSample Sentence:\n=====')
      print(sample_sentence(hmm4, obs_map, n_words=25))
```

Sample Sentence:

=====

Not cases have of judicial blood impeachment law their the fourths or be shall
of states of the place adjournment concurrence of representatives or on...

```
[9]: hmm16 = unsupervised_HMM(obs, 16, 100)
print('\nSample Sentence:\n=====')
print(sample_sentence(hmm16, obs_map, n_words=25))
```

Sample Sentence:

=====

To inferior of a monday maritime the judges of behaviour and arsenals and taxes
been yeas concur departments and diminished to work different may a...


1.6 Part I: Visualizing the wordcloud of each state

Below, we visualize each state as a wordcloud by sampling a large emission from the state:

```
[10]: wordclouds = states_to_wordclouds(hmm8, obs_map)
```



State 1



State 2



State 3



State 4



State 5



State 6




State 7



State 8



State 9



1.7 Visualizing the process of an HMM generating an emission

The visualization below shows how an HMM generates an emission. Each state is shown as a wordcloud on the plot, and transition probabilities between the states are shown as arrows. The darker an arrow, the higher the transition probability.

At every frame, a transition is taken and an observation is emitted from the new state. A red arrow indicates that the transition was just taken. If a transition stays at the same state, it is represented as an arrowhead on top of that state.

Use fullscreen for a better view of the process.

```
[11]: anim = animate_emission(hmm8, obs_map, M=8)
      HTML(anim.to_html5_video())
```

Animating...

```
[11]: <IPython.core.display.HTML object>
```

